

CAB420: Dimension Reduction

DIMENSIONS ARE BAD, MMMKAY

Why Reduce Dimensions?

- True vs Observed Dimensionality
 - Consider 20x20 binary bitmaps of handwritten digits
 - We have $\{0,1\}^{400}$ possible patterns
 - Most of these will never be seen
 - True dimensionality
 - Possible variations of the pen stroke
 - The set of examples we can actually see




Curse of Dimensionality

Machine learning methods are based on statistics

- More observations means more reliable statistics

But

- More dimensions (without more examples) means that our examples are sparser in our feature space
- Less reliable models



Why Reduce Dimensions?

- Simplify or reduce the data
 - Help better represent the data given a limited number of samples
- Not all dimensions are equal
 - Some are high informative, others are not so useful, it's good to get rid of the useless ones
- Can make other machine learning tasks easier
 - And faster


Principal Component Analysis

PCA

Principal Component Analysis


Principal components are:

- The directions where there is the most variance
- Can be seen as the underlying structure of the data



Eigenvectors and values

- Exist in pairs
 - Eigenvectors are the directions of variance
 - Eigenvalues specify the amount of variance in for the eigenvector
- We have as many pairs as we do dimensions in the source data



Finding Eigenvalues and Eigenvectors

- A covariance matrix can be formed between all dimensions in a dataset
- This matrix will be the same whether or not our data is centered (i.e., every entry in a dimensions is reduced by its mean)
- Principal components (PCs) are the directions in which variance is changing the most, not typically along an axis line
- PCs are represented by both eigenvalues and eigenvectors
 - Eigenvalues are related to the amount of variance in a principal component
 - Eigenvectors determine the direction of the PCs variance
 - These eigenvectors must have a magnitude of 1, as they will become unit measures.
- Recovery of Eigenvalues and Eigenvectors is done via linear algebra
 - Outside the scope of this subject
- Can also be found using Singular Value Decomposition
 - Typically, more numerically stable

Principal Component Analysis

- As the output of our PCA process, we get

$$T = XW$$

- X is the input data, containing N samples and P dimensions
- T is the transformed data
 - Same size as X
- W is the learned transformation
 - $P \times P$ matrix
 - Each column describes how the P dimensions are combined to create a new dimension
- To compute PCA, we need $N > P$
 - More samples than dimensions

Principal Component Analysis

- PCA will also offer a measure of variance for each new dimension
 - Derived from the Eigenvalues
 - New dimensions are returned in order of variance
 - The first few dimensions can be seen as the most important
 - Contain the most variance, thus the most information
 - Dimensions that contain limited variance are closer to being a constant, thus contain less information
- Using data that is not standardised can distort PCA results

Principal Component Analysis

Doesn't

- Add new dimensions
- Remove dimensions
- Change the data

It does

- Re-project the data along a new of axes such that
 - The first dimension has the most variance
 - The second dimension has the second most variance
 - And so on


Projecting the Data

- When projecting it also centres the data
 - Translations of the data have no impact on the PCA results
- Projected axes are orthogonal to each other
 - i.e. at right angles
 - Ensures that the new axes can cover the data space as efficiently as possible
- We can also transform from the PCA space back to the original space

$$\hat{X} = TW'$$


What are the new Dimensions?

- The new dimensions can be seen as weighted sums of the existing ones
 - Our first PC is (roughly)
 - $0.5x + 0.5y$
 - Our second PC is (roughly)
 - $-0.5x + 0.5y$
- PCA can be seen as computing a rotation matrix and rotating the data




Simple Example

- See ***CAB420_Dimension_Reduction_Example_1_Principal_Component_Analysis.ipynb***
- Our data
 - Some random points on a line
- After applying PCA
 - Points are "lined up" on X and Y dimension
 - Centred
 - The way the points are distributed (their shape) is the same




Simple Example: Reprojection

- When we reproject the data, we get a shifted version
 - PCA centers the data for us
 - Set the mean to be 0
 - First step of standardisation
 - Our reprojected version is thus centered at the origin
 - We need to save the mean of the original data to totally recover the original
 - Python will save the mean for us



Simple Example

- Projection matrix, W , shows how new dimensions are created
 - New dimensions are a weighted combination of original dimensions
- Eigenvalues show variance in each new dimension
 - First dimension contains 99.5% of the variance
 - This makes sense, our original data is spread out along one narrow line



$$W = \begin{bmatrix} -0.260 & -0.966 \\ -0.966 & 0.260 \end{bmatrix}$$

$$\text{Eigenvalues} = \begin{bmatrix} 12.144 \\ 0.050 \end{bmatrix}$$

CAB420: PCA for Dimension Reduction

WE'VE GOT SOME NEW AXES, NOW WHAT?

Reducing Dimensions

PCA gives us

- A new coordinate frame
- Knowledge about how much of the total variance is captured by each new dimension

So we can just select the top N dimensions

- Retain a given amount of variance
- Remove dimensions that contribute very little

Reducing Dimensions

- We can view this as being

$$T_L = XW_L$$

- L is the number of dimensions we wish to keep
 - We select the first L columns of W , W_L
 - Our output now has only L dimensions

- We can reproject from the reduced space using

$$\hat{X} = T_L W_L'$$

- Information will be lost in such a reprojection

An Example

- See ***CAB420_Dimension_Reduction_Example_2_PCA_and_Dimension_Reduction.ipynb***
- Our data:
 - The Iris dataset, measurements of Iris petals, 4 dimensions
- Approach:
 - Apply PCA and transform data
 - Reconstruct the original using only some of the dimensions

Computed Transform

- When we compute PCA, we get
 - A transform
 - We have 4 dimensions in our input, so our transform matrix is 4x4
 - A measure of variance in each new dimension
 - The explained variance ratio is shown below (rather than the raw Eigenvalues)

```
[ [ 0.361 -0.085  0.857  0.358 ]  
  [ 0.657  0.730 -0.173 -0.075]  
  [-0.582  0.598  0.076  0.546]  
  [-0.315  0.320  0.480 -0.754]]
```

```
[0.92461872  
 0.05306648  
 0.01710261  
 0.00521218]
```

1 Dimension

- Reconstructed and actual values for first 5 samples
 - Reconstructed samples are close to the original in values
 - We had ~92.5% of variance in the first principal component

Reconstructed:

```
[ [4.86247892  3.28673941  1.4328746   0.22688569]
  [4.79929088  3.30151808  1.2830867   0.16423922]
  [4.85120324  3.28937661  1.40614547  0.21570665]
  [4.85721176  3.28797132  1.42038875  0.22166368]
  [5.01906124  3.25011732  1.8040546   0.38212597] ]
```

Actual:

```
[ [4.9  3.   1.4  0.2]
  [4.7  3.2  1.3  0.2]
  [4.6  3.1  1.5  0.2]
  [5.   3.6  1.4  0.2]
  [5.4  3.9  1.7  0.4] ]
```

Error: 0.085604

2 Dimensions

- Reconstructed and actual values for first 5 samples
 - Slightly more accurate than 1 dimension

Reconstructed:

```
[[4.7462619  3.15749994 1.46356177 0.24024592]  
 [4.70411871 3.1956816  1.30821697 0.17518015]  
 [4.6422117  3.05696697 1.46132981 0.23973218]  
 [5.07175511 3.52655486 1.36373845 0.19699991]  
 [5.50581049 3.79140823 1.67552816 0.32616959]]
```

Actual:

```
[[4.9 3.  1.4 0.2]  
 [4.7 3.2 1.3 0.2]  
 [4.6 3.1 1.5 0.2]  
 [5.  3.6 1.4 0.2]  
 [5.4 3.9 1.7 0.4]]
```

Error: 0.025341

3 Dimensions

- Reconstructed and actual values for first 5 samples
 - Minimal error with three dimensions included

Reconstructed:

```
[[4.86875839 3.03166108 1.4475168 0.12536791]
 [4.69370023 3.20638436 1.30958161 0.18495067]
 [4.6238432 3.07583667 1.46373578 0.25695828]
 [5.0193263 3.58041421 1.37060574 0.24616799]
 [5.40763506 3.89226243 1.68838749 0.41823916]]
```

Actual:

```
[[4.9 3. 1.4 0.2]
 [4.7 3.2 1.3 0.2]
 [4.6 3.1 1.5 0.2]
 [5. 3.6 1.4 0.2]
 [5.4 3.9 1.7 0.4]]
```

Error: 0.005919

Breaking PCA

BREAKING STUFF IS FUN

PCA


- Dimension reduction technique
 - Project the data into a set of new dimensions such that the
 - First new dimension captures the most variance
 - The second captures the second most variance
 -
- Can reduce dimensionality by electing to only retain a percentage of total variance
 - Typically most (90% or more) variance is in a small number of dimensions (10% or less)

But...

- Is the amount of variance captured the best criteria for determining what to keep?
- Are the most important variables the ones that make different types of object distinct?

Breaking PCA

- See *CAB420_Dimension_Reduction_Example_3_Linear_Discriminant_Analysis.ipynb* (the first part)
- Simple dataset
 - Two classes
 - Well separated



Breaking PCA


- Computing PCA on our dataset, we get:
 - Explained variance: [0.85471369 0.14528631]
 - Most variation is in the first dimension
- If we transform the data using just the first PC, and reproject we get:

Actual:	Reconstructed:
[3 4]	[3.73842482 3.44234505]
[4 5]	[4.58253002 4.56007607]
[5 6]	[5.42663522 5.67780709]
[5 7]	[5.90755333 6.31462]
[2 1]	[1.9324834 1.05098821]
[3 2]	[2.7765886 2.16871923]
[4 2]	[3.13977569 2.64963734]
[4 3]	[3.6206938 3.28645025]
[6 4]]	[4.82798609 4.88509938]]

- Average reconstruction error = 0.408243


Breaking PCA

- Left: original data
- Right: reconstructed data (from first PC)
- We've lost any ability to separate the classes



Breaking PCA

- If we used only the second principal component
 - Our reconstruction looks nothing like the original data
 - Class separation is preserved



CAB420: Linear Discriminant Analysis

A BIT LIKE PCA, BUT NOT

Linear Discriminant Analysis

- Also often called Fisher Discriminant Analysis
- LDA seeks to
 - Find a projection that maximises the ratio between the inter class and between class scatter matrices
 - Meaning:
 - Samples in the same class get closer together
 - Samples in different classes get further away
 - Find a projection that best **discriminates** between the classes

LDA Overview

- PCA considers the variance between existing dimensions of the data
 - Finds a new set of dimensions such that the first dimension contains the most variance, etc.
- In Linear Discriminant Analysis, we instead focus on the "variance" relating to classes of the data.
- In order to effectively separate clusters during dimension reduction, we must take into consideration:
 1. The "variance" between classes, and
 2. The "variance" within each class.

LDA Overview

- The difference between each point and its cluster's mean is the crucial part of LDA.
 - We don't need the "divided by $n - 1$ " that we normally associate with variance. In fact, this will only scale our new dimensions.
 - The matrix result without the division is instead called a **scatter matrix**.
- One important note is that we should ensure a sufficient number of points in each class in order for LDA to succeed. This is because LDA is very sensitive to outliers.
- The number of data points in the smallest class should be larger than the number of explanatory variables.

LDA – Scatter Matrices

- LDA is based on scatter matrices
- Within class
 - $S_w = \sum_{i=1}^n (x_i - \mu_{y_i})(x_i - \mu_{y_i})^T$
 - where
 - n is the number of points
 - μ_{y_i} is the mean of the i th class
- Between class
 - $S_B = \sum_{k=1}^m n_k (\mu_k - \mu)(\mu_k - \mu)^T$
 - where
 - m is the number of classes
 - n_k is the number of points in the k th class
 - μ_k is the mean of the k th class
 - μ is the mean of the data


Scatter Matrices Visualised

- See ***CAB420_Dimension_Reduction_Example_3_Linear_Discriminant_Analysis.ipynb*** (the second part)
- 2D Data, 3 Classes
 - Within class scatter

```
[[0.01331477 0.0015448 ]  
 [0.0015448  0.01423041]]
```

- Between class scatter

```
[[248.86709419 249.70285217]  
 [249.70285217 250.54892153]]
```



Scatter Matrices Visualised


- 2D Data, 3 Classes

- Within class scatter

```
[[ 0.35319307 -0.03366045]  
 [-0.03366045  0.59361245]]
```

- Between class scatter

```
[[266.73188433 251.55160505]  
 [251.55160505 237.23710806]]
```



Scatter Matrices Visualised


- 2D Data, 3 Classes

- Within class scatter

```
[[0.74859984 0.41766819]  
 [0.41766819 0.89271018]]
```

- Between class scatter

```
[[ 0.22332289 -0.19700407]  
 [-0.19700407 0.17936888]]
```



Computing LDA

- LDA seeks to maximise ratio of between class to within class scatter

$$\hat{w} = \operatorname{argmax}_w \frac{w^T S_B w}{w^T S_W w}$$

- This can be solved by computing the eigenvectors for $\frac{S_B}{S_W}$.
- LDA will only return $C - 1$ meaningful components
 - C is the number of classes in the data

Computing LDA

- Same simple sample data we broke PCA with (on the left)

- Within class scatter matrix

```
[[24.13333333 24.      ]  
 [24.         26.      ]]
```

- Between class scatter matrix

```
[[ 0.77575758 -2.90909091]  
 [-2.90909091 10.90909091]]
```

- Transformed Data (on the right)



- Eigenvectors

```
[[-0.96623494  0.71169327]  
 [-0.25766265 -0.70249034]]
```

- Eigenvalues

```
[0.         8.22044277]
```

- Only 1 non-zero Eigenvalue

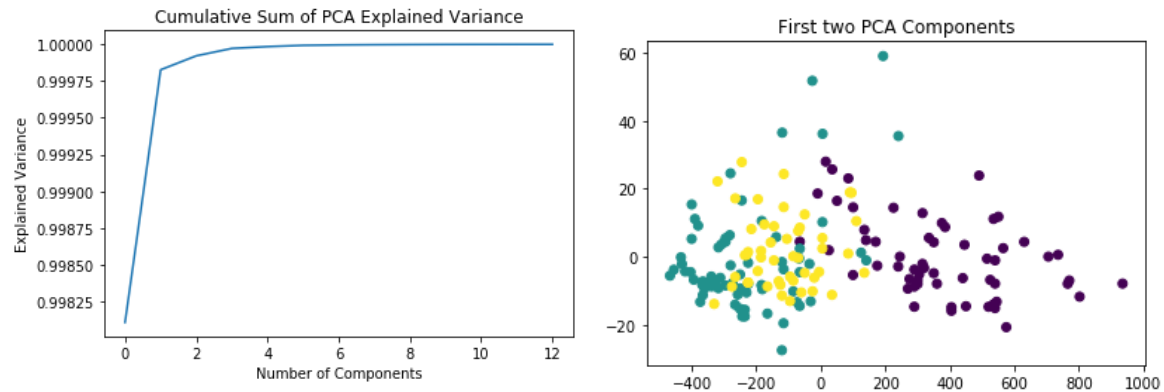


An Example – PCA vs LDA

- See *CAB420_Dimension_Reduction_Example_4_Linear_Discriminant_Analysis_II_Action_Time.ipynb*
- Data
 - Chemical properties of wine from three different cultivars in Italy
 - 12 variables
 - This is ostensibly a classification task, though we'll just analyse the data and visually look at class separation



PCA

- Almost all the information is in the first dimension
 - 99.8% of the variance in one dimension
- Plotting the top two dimensions (~99.975% of the variance) we see that we have some class separation
 - Purple is somewhat separated, the other classes not so much



LDA

- Note that I only have two meaningful dimensions
 - LDA returns $C - 1$ components
 - I have 3 classes, thus 2 components
- Two dimensions result in very good separation of classes




Breaking LDA

- LDA is sensitive to the number of samples per class
- If we have very few samples per class
 - Scatter matrices become hard to predict
 - LDA solution can become unstable
- This problem becomes more pronounced with different solvers
 - SVD (Python's default) is generally more stable and robust in these situations

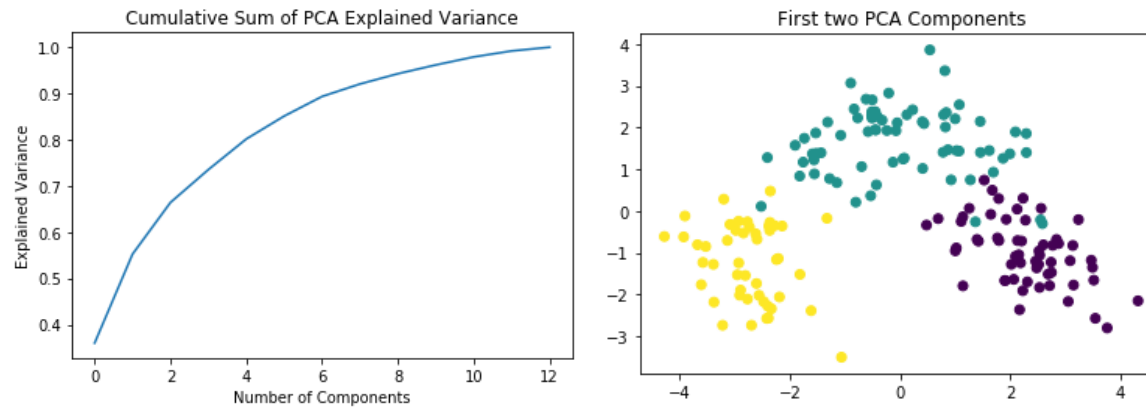
PCA – what went wrong?

- The boxplot of the data explains our problem
 - The variation in the data is dominated by a single dimension
 - This distrorts our PCA, and hurts it's performance



Improving PCA

- Standardised data
- We no longer have all the variation in one dimension
- Class separation much better on first two dimensions
 - Not as good as LDA, but close



LDA and Acronym Overuse

Warning

- There is another LDA: latent Dirichlet allocation
- Some toolkits/text refer to latent Dirichlet allocation as LDA exclusively
 - If you search for LDA in MATLAB you will get taken to info on latent Dirichlet allocation
 - Google will give you results for both
- “Fisher Discriminant Analysis” is another name for Linear Discriminant Analysis

Incidentally...

- latent Dirichlet allocation is a really cool method for modelling topic distributions
- We’ll look at that when we look at sequences

CAB420: t-SNE

YET MORE DIMENSION REDUCTION,
YET ANOTHER ACRONYM


t-distributed Stochastic Neighbour Embeddings

t-SNE is a bit different

- Fairly recent method (2008) developed for visualisation
- Not a linear projection
- Projects data into two dimensions
- Uses local relationships to create a low dimensional mapping that can capture non-linear structure

Why t-SNE?

- Lots of things aren't linear (see left), and PCA (and LDA) struggle to capture such a relationship
- Overcomes the “crowding problem”, when projecting from a high to low dimensional space points can become tightly clustered and “crowded”



Why t-SNE?

- Intended for visualisation
 - Often we use dimension reduction to visually analyse data, t-SNE is designed for this
- Becoming widely used
 - Used a lot when visualising neural network intermediate states or outputs
 - We will use it for this and similar purposes
 - Important to understand how it differs from other dimension reduction techniques

t-SNE Overview

Two main steps:

- Step 1: In high dimensional space, create a probability distribution that captures the relationship between points
- Step 2: Create a low dimensional space that follows the created distribution as best as possible
- More details: <https://mlexplained.com/2018/09/14/paper-dissected-visualizing-data-using-t-sne-explained/>
- <https://distill.pub/2016/misread-tsne/>

Limitations

- Non-convex optimisation
 - Not guaranteed to reach a global minima (i.e. a constant best solution)
 - Non-deterministic
 - If you run it again, you'll get different results
- Assumes Locally Linear Manifolds
 - Uses Euclidean distance to compare points (which is generally an ok assumption)
 - If local relationships are highly non-linear, may struggle to make sense of things

CAB420: Eigenfaces

FACE REC WITH PCA

Face Recognition: Eigenfaces

- Classic face recognition approach that uses PCA
- See ***CAB420_Dimension_Reduction_Example_5_Eigenfaces.ipynb***
- Problem
 - Identify a person from an image, given an input image and a database of images with known identities

Face Recognition: Naïve Solution


- For input image
 - Align/crop image such that it's at a consistent pose/scale
 - Compare, pixel-wise, to all images in the dataset
 - Select identity of closest image as best match
- Problems with this approach
 - Image are large, even small images have very high dimensionality
 - What happens in the lighting changes, the background changes, etc?

Eigenfaces

- For our database
 - Compute PCA over database, return top N dimensions
- For input image
 - Map to PCA space
 - Compare to database in PCA space and select best match as target ID
- Still sensitive to lighting, background, etc
 - Though less so than a raw pixel approach
- Much quicker due to fewer retained dimensions

Input Data

- Face images
 - 32 x 32 pixels, greyscale
 - Varied lighting
 - Consistent pose
 - Images normalised based on eye locations



PCA and Faces

- The mean face
 - Fairly average looking
 - Perhaps more male than female
 - Nature of the dataset, more males than females in the data
 - All lighting variation removed



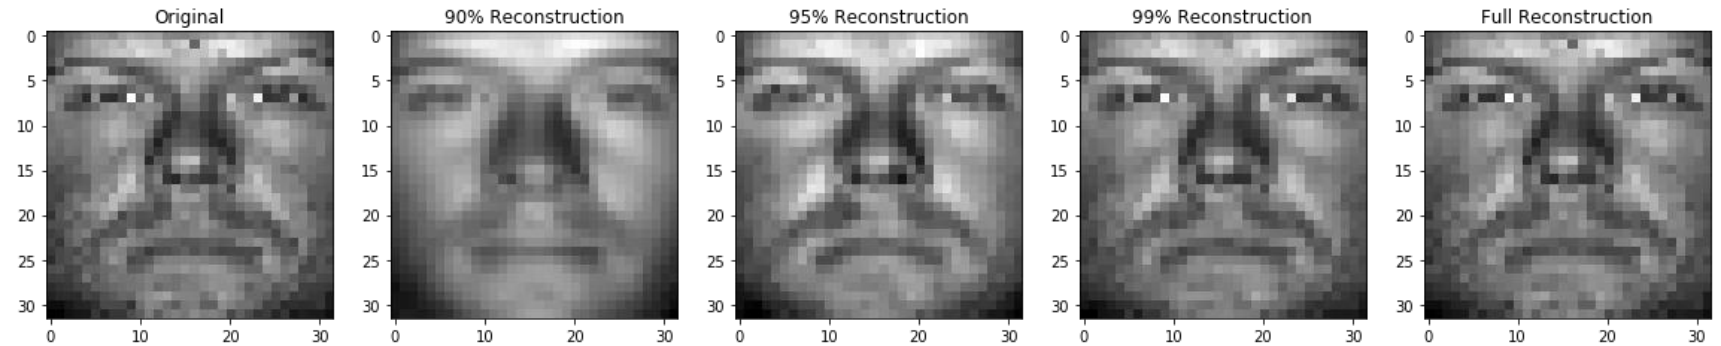
Eigenfaces

- Our data are images
 - Thus, our eigenvectors (principal components) are images
- Our eigenvectors (eigenfaces) capture the major variations in faces
 - Early eigenvectors capture the drastic lighting changes





Low Dimensional Faces

- Reconstructions as retained PCs are increased
 - 90% is 23 PCs
 - 95% is 60 PCs
 - 99% is 238 PCs
- More PCs means more fine-grained details
 - Fewer PCs can be (sort-of) seen as a low pass filter



Recognising Low-Res Faces

- Train a classifier
 - CKNN
 - Trained on dimension reduced data



Combining PCA and LDA

LIVING TOGETHER, IN PERFECT HARMONY
(OR SOMETHING LIKE THAT)

Why?

Don't they do similar things?

- Reduce dimensions, etc?

Sort of, but not really

- PCA is focussed on reconstruction
- LDA focussed on class separation
 - Throws data away

LDA can have problems when you have

- Lots of dimensions and/or classes
- Few sample points per class

When LDA Goes Bad

When we perform LDA with

- Lots of dimensions and/or classes; and/or
- Few sample per class

The estimate of the within scatter class matrix in particular become poor

- Can lead to overfitting
- Can lead to a singular or close to singular matrix
 - Singular \rightarrow non invertible
 - We need to invert with the within class scatter matrix
- Leads to bad results

PCA as Pre-Processing

A Solution

- Reduce the size of the data space
- i.e. PCA

The process

- Apply PCA, select top N components, removing uninformative dimensions
- Apply LDA to reduced space
- When transforming a new sample point:
 - $x_{\text{transformed}} = x * \text{PCA} * \text{LDA}$
 - i.e. transform to PCA space, then to LDA space

This will be explored further in this weeks tutorial

Final Thoughts

SUMMING UP AND ALL THAT

LDA vs PCA

- PCA
 - Unsupervised
 - Aim to preserve as much information as possible
 - Can be inverted
- LDA
 - Supervised
 - We need to know class labels
 - Aim to preserve as much discriminative power as possible
 - Cannot be inverted

PCA Requirements

- PCA is unsupervised
 - No labels needed, just the raw data
- We need to have more samples than we do dimensions
 - If we don't have this, many PCA methods will transpose the data
 - Or be unstable
- Standardisation can help (a lot)

PCA Requirements

- Ideally, we want many more samples than dimensions
 - PCA may become unstable if we only have slightly more samples than dimensions
- We can do robustness tests by computing multiple PCA transforms on datasets sampled from our overall data
 - If the PCA transform is consistent, we can claim that we have enough data
 - Similar to the idea of the standard error measured when fitting a regression model

LDA Requirements

- Class labels
 - LDA is supervised
- Sufficient examples per class to estimate scatter matrices
 - Ideally, more samples per class than dimensions
- A problem that doesn't require reconstruction
 - We cannot reconstruct the original signal from LDA
- Standardisation not needed
 - The scatter matrices capture this
 - Standardisation may result in an axis being flipped, but that's it

Which one to use?

As the previous slides suggests, it depends on

- What are you trying to do?
 - Classification?
 - Visualisation?
 - Compression?
- What data do you have?
 - Class labels or not?
- Methods can be used in combination
 - FisherFaces: apply PCA, then apply LDA

Lots of other dimension reduction techniques too

The methods that we've look at are not always best. Other methods include:

- Independent Component Analysis
- Factor Analysis (and it's many variants)
- Probabilistic PCA and Probabilistic LDA
- Deep network-based methods
 - Metric learning
 - Auto-encoders
- And many, many, more ...